



**University of  
Zurich<sup>UZH</sup>**



MAKERERE UNIVERSITY

# **Data Analysis with R:**

## **Lecture Slides (all)**

Sonja Hartnack, Terence Odoch & Muriel Buri

October 2018

# Goals of the course

To be able to...

- import data sets to R
- describe data with R
- apply basic statistical tests in R
- some ideas for more advanced statistical tools ...
- simulate a data set similar to own research

# Getting to know each other

- My name is ...
- I am doing a Master / a PhD in ...
- I hope to learn in this course how to ....
- My personal goal for this course is ...

# How do we reach these goals

- hands on exercises with R:
  - chickwts
  - ToothGrowth
  - bacteria
  - perulung
  - ... and others.
- interactive discussions
- asking a lot of questions: ask google!
- group work
- short motivational lectures

## Get started with data set: chickwts

An experiment was conducted to measure and compare the effectiveness of various feed supplements on the growth rate of chickens.


```
data(chickwts)
head(chickwts)
```

```
##   weight      feed
## 1    179 horsebean
## 2    160 horsebean
## 3    136 horsebean
## 4    227 horsebean
## 5    217 horsebean
## 6    168 horsebean
```

# Ideas for plotting the data

# Ideas for plotting the data

```
boxplot(weight ~ feed, data = chickwts)
```



# Ideas for analysing the data



# Ideas for analysing the data

```
summary(aov(weight ~ feed, data = chickwts))
```

```
##              Df Sum Sq Mean Sq F value    Pr(>F)
## feed           5 231129   46226    15.37 5.94e-10 ***
## Residuals     65 195556     3009
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

# Functionalities in R and RStudio



- Source
- Console
- Environment, History, Files
- Files, Plots, Packages, Help

The screenshot displays the RStudio environment with the following components:

- Environment Pane:** Shows the Global Environment with variables: `ccatable` (num [1:3, 1:2] 62 16 22 163 20 17), `francrostudy` (1410 obs. of 13 variables), `mirtable` (num [1:2, 1:2] 10 26 11027 11008), `dielttable` (num [1:2, 1:2] 74 61 44 2383 3747 ...), `franz` (1410 obs. of 13 variables), and `salarytable` (num [1:2, 1:2] 122 64 75 50).
- Console:** Displays the R startup message, including the license and contributors information, and the workspace load status: `[Workspace loaded from ~/.RData]`.
- Files Pane:** Shows a file explorer view of the Home directory, listing files such as `JfData` (42.2 KB), `JfHistory` (13.1 KB), `SoftChalDownloadTempDir`, `Adabo`, `Backup Jump Drive`, `Buildings study 12-9-16.rgs` (1.2 MB), `Centasis Studio`, `Confounding and Effect Measure Modification_JL.docx` (21.1 KB), `Confounding Workshop Explanation.xlsx` (11.4 KB), `Custom Office Templates`, `i-Random linear pptx` (2.6 MB), `FP713`, and `Humeral Response.txt` (882 B).



- Define manually a new folder called **RCourse** in your personal documents on your personal computer
- Know in which directory you are

```
getwd()
```

```
## [1] "/home/mburi/ownCloud/git/DataAnalysisWithR/Lectures"
```

- Set directory path

```
# back- and forslah is dependent on the system
```

```
setwd("C:/Users/mburi/Documents/RCourse/")
```

```
setwd("C:\\Users\\mburi\\Documents\\RCourse\\")
```

- Always clean up before starting with new R-Script

```
rm(list=ls()) # empty workspace, delete previously saved variables
```



```
?chickwts
```

Also, have a look at the examples at the end of the help pages.

# Exercise 1



## A data frame in R: chickwts



chickwts[ ROWS , COLUMNS ]

	weight	feed
1	179	horsebean
2	160	horsebean
3	136	horsebean
4	227	horsebean
5	217	horsebean
6	168	horsebean
7	108	horsebean
8	124	horsebean
9	143	horsebean
10	140	horsebean
11	309	linseed
12	229	linseed
13	181	linseed

chickwts[ 6, 1 ]

	weight	feed
1	179	horsebean
2	160	horsebean
3	136	horsebean
4	227	horsebean
5	217	horsebean
6	168	horsebean
7	108	horsebean
8	124	horsebean
9	143	horsebean
10	140	horsebean
11	309	linseed
12	229	linseed
13	181	linseed

chickwts[ 11, 2 ]



### Values of ...

```
# ... all columns of sixth observation:  
chickwts[6, ]  
  
# ... all columns of sixth to eleventh observation:  
chickwts[6:11, ]  
  
# ... all columns of sixth, eleventh and twentieth observation:  
chickwts[c(6, 11, 20), ]  
  
# ... all rows of first column (weight):  
chickwts[, 1]  
  
# ... all rows of second column (feed):  
chickwts[, 2]  
# or use the "$" sign as a reference to column "feed":  
chickwts$feed
```

## What is a data frame in R?



A data frame is used for storing a list of vectors of equal length. For example, the following variable `df` is a data frame containing three vectors `n`, `s`, `b`.

```
n <- c(2, 3, 5)
s <- c("aa", "bb", "cc")
b <- c(TRUE, FALSE, TRUE)
df <- data.frame(n, s, b) # df is a data frame
```

Following are the characteristics of a data frame:

- The column names should be non-empty.
- The row names should be unique.
- The data stored in a data frame can be of numeric, factor or character type.
- Each column should contain same number of data items.





```
a <- c(1, 2, 3, 4)
```

```
a
```

```
## [1] 1 2 3 4
```

```
data.frame(a)
```

```
##      a
```

```
## 1 1
```

```
## 2 2
```

```
## 3 3
```

```
## 4 4
```

```
b <- c("d", "h", "h", "d")
```

```
dat <- data.frame(a, b)
```

```
dat
```

```
##      a b
```

```
## 1 1 d
```

```
## 2 2 h
```

```
## 3 3 h
```

```
## 4 4 d
```

# Data frame in R: How to add a vector variable



```
c <- factor(c("male", "female", "male", "female"),
            levels = c("female", "male"))
dat$c <- c # use "$" to refer to the additional vector variable
dat
```

```
##   a b      c
## 1 1 d   male
## 2 2 h female
## 3 3 h   male
## 4 4 d female
```

```
# What is the dimension (number of rows and columns) of our data frame?
dim(dat) # 4 rows and 3 columns
```

```
## [1] 4 3
```

## Exercise 2



# RELATIONSHIPS IN R WOLRD - What is a package within R



```
# install.packages("MASS")  
# library("MASS")  
# data(bacteria)  
# ?bacteria
```

# WHAT IS A PACKAGE AND How to install IT



```
install.packages("MASS")  
library("MASS")  
data(bacteria)  
?bacteria
```

## Exercise 3



# How to google for getting help in R

- Google for **select observations in R**.



Objects are assigned values using  $<-$ , an arrow formed out of  $<$  and  $-$ . For example, the following command assigns the value 1 to the object a.

```
a <- 1 # ALWAYS use "gets" assignment operator!  
# a = 1 # DO NOT USE the equal sign as the assignment operator!
```

After this assignment, the object a contains the value 1. Another assignment to the same object will change the content.

```
a <- 5
```



## Examples of assigned objects: Single number



```
a <- 1
b <- 2
c <- a + b # c = 3
c

## [1] 3
```

# Examples of assigned objects: Vector



```
a <- c(1, 2, 3, 4, 5)
```

```
b <- 1
```

```
c <- a + b
```

```
c
```

```
## [1] 2 3 4 5 6
```

## Examples of assigned objects: Model



```
anova_model <- aov(weight ~ feed, data = chickwts)
summary(anova_model)
```

```
##              Df Sum Sq Mean Sq F value    Pr(>F)
## feed           5 231129   46226    15.37 5.94e-10 ***
## Residuals     65 195556     3009
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

# Examples of assigned objects: Data frame



```
bac <- bacteria
str(bac) # $ week: int  0 2 4 11 0 2 6 11 0 2 ...

## 'data.frame': 220 obs. of  6 variables:
## $ y   : Factor w/ 2 levels "n","y": 2 2 2 2 2 2 1 2 2 2 ...
## $ ap  : Factor w/ 2 levels "a","p": 2 2 2 2 1 1 1 1 1 1 ...
## $ hilo: Factor w/ 2 levels "hi","lo": 1 1 1 1 1 1 1 1 2 2 ...
## $ week: int  0 2 4 11 0 2 6 11 0 2 ...
## $ ID  : Factor w/ 50 levels "X01","X02","X03",...: 1 1 1 1 2 2 2 2 3 3 ...
## $ trt : Factor w/ 3 levels "placebo","drug",...: 1 1 1 1 3 3 3 3 2 2 ...

bac_sub <- subset(bac, week == 2)
str(bac_sub) # $ week: int  2 2 2 2 2 2 2 2 2 2 ...

## 'data.frame': 44 obs. of  6 variables:
## $ y   : Factor w/ 2 levels "n","y": 2 2 2 2 2 2 1 2 2 2 ...
## $ ap  : Factor w/ 2 levels "a","p": 2 1 1 2 2 1 1 2 2 2 ...
## $ hilo: Factor w/ 2 levels "hi","lo": 1 1 2 2 2 2 1 1 2 1 ...
## $ week: int  2 2 2 2 2 2 2 2 2 2 ...
## $ ID  : Factor w/ 50 levels "X01","X02","X03",...: 1 2 3 4 5 6 7 8 9 11 ...
## $ trt : Factor w/ 3 levels "placebo","drug",...: 1 3 2 1 1 2 3 1 1 1 ...
```



The `str` function displays the structure of an R object. One line for each "basic" structure is displayed.

```
## 'data.frame': 44 obs. of 6 variables:
## $ y : Factor w/ 2 levels "n","y": 2 2 2 2 2 2 1 2 2 2 ...
## $ ap : Factor w/ 2 levels "a","p": 2 1 1 2 2 1 1 2 2 2 ...
## $ hilo: Factor w/ 2 levels "hi","lo": 1 1 2 2 2 2 1 1 2 1 ...
## $ week: int 2 2 2 2 2 2 2 2 2 2 ...
## $ ID : Factor w/ 50 levels "X01","X02","X03",...: 1 2 3 4 5 6 7 8 9 11 ...
## $ trt : Factor w/ 3 levels "placebo","drug",...: 1 3 2 1 1 2 3 1 1 1 ...
```

## Exercise 4



# Data types in R



- numeric

```
data(ToothGrowth)
ToothGrowth$len[1:6]

## [1]  4.2 11.5  7.3  5.8  6.4 10.0

class(ToothGrowth$len[1:6])

## [1] "numeric"
```

- integers

```
bacteria$week[1:6]

## [1]  0  2  4 11  0  2

class(bacteria$week[1:6])

## [1] "integer"
```

- (un/ordered) factor

```
chickwts$feed[1:6]

## [1] horsebean horsebean horsebean horsebean horsebean horsebean
## Levels: casein horsebean linseed meatmeal soybean sunflower

levels(chickwts$feed)[1:3]

## [1] "casein"      "horsebean"  "linseed"
```



Ordinal variables are represented as ordered factors:

```
bac_growth <- c("none", "+", "++", "+", "+++", "+", "none") # vector
bac_growth <- factor(bac_growth, levels = c("none", "+", "++", "+++"),
                     order = TRUE)

bac_growth

## [1] none +      ++      +      +++     +      none
## Levels: none < + < ++ < +++

#
mood <- c("OK", "Well", "Super", "Super", "Don't ask", "OK") # vector
mood <- factor(mood, levels = c("Don't ask", "Well", "OK", "Super"),
               order = TRUE)

mood

## [1] OK          Well          Super          Super          Don't ask OK
## Levels: Don't ask < Well < OK < Super
```



## Exercise 5



## Exercise 6



## Rules for importing data into R (from Excel)



- First row of excel sheet contains **variable names**:  
y, ap, hilo, week, ID, trt.
- Columns of excel sheet represent **variables**.
- Rows of excel sheet represent **observations per individual** (except for the first row).

	A	B	C	D	E	F	G	H	I	J
1	id	fev1	age	height	sex	respsymptoms				
2	1	1.56	9.593	124.8	0	0				
3	2	1.18	7.491	111	1	0				
4	3	1.87	9.864	135.7	0	0				
5	4	1.49	8.588	119.1	0	0				
6	5	1.62	8.967	120.9	1	0				
7	6	2.11	9.293	134.3	0	1				
8	7	1.73	9.574	122.1	1	0				
9	8	1.47	8.493	122.6	0	1				
10	9	1.83	8.468	126.8	1	0				
11	10	1.41	9.029	126	0	0				
12	11	1.27	8.274	128	0	0				
13	12	1.34	8.416	127	0	0				
14	13	1.64	9.629	133.7	0	0				
15	14	1.57	8.622	125.5	1	0				
16	15	1.51	9.033	125.9	1	0				
17	16	1.25	8.643	122.3	0	0				



Variable names should ...

- start with a letter (not a number):  
y, ap, hilo, week, ID, trt
- longer variables names should be separated with dots:  
`time.in.weeks`
- do not use special characters, such as /, #, @, &, \*, ...

# How to import external data files into R?



- Excel Files (.xlsx)

```
install.packages("readxl")  
library("readxl")
```

- CSV Files (.csv)

```
?read.csv
```

- Text files (.txt)

# How to import Excel files into R?

## Three major steps: Excel file preparation



# How to import Excel files into R?

Three major steps: Save raw data file as .csv



# How to import Excel files into R?

## Three major steps: Import data into R





# How to import Excel files into R?

## Three major steps: Import data into R



# How to import Excel files into R?

## Three major steps: Import data into R



# How to import Excel files into R?



```
# Import .csv file with the help of the read.csv function  
# Be sure to add sep = ";" so that we separate the columns.  
lung <- read.csv("C:\\Users\\Exercises\\data\\perulung_ems.csv",  
                 sep = ";")  
  
head(lung)  
str(lung)
```



Data from a study of lung function among children living in a deprived suburb of Lima, Peru. Data taken from Kirkwood and Sterne, 2nd edition.

Variables:

- fev1: in liter, "forced expiratory volume in 1 second" measured by a spirometer. This is the maximum volume of air which the children could breath out in 1 second
- age: in years
- height: in cm
- sex: 0 = girl, 1 = boy
- respsymp: respiratory symptoms experienced by the child over the previous 12 months

# Why do we need Statistics?

## **Repeatability of results:**

**Statistical science** allows us to estimate what might happen if an experiment was repeated - but without having to actually repeat it!

## Why do we need Statistics?

- Study results must be shown to be robust, i.e. real and not due to random chance
- Best way to demonstrate this is to repeat the same experiment/study many times each with **different** subjects (animals) drawn from the **same study population** and show that the result is truly repeatable
- It is generally totally impractical, in terms of both time and resources, to repeat an experiment many times!

# Why do we need Statistics?

- Instead of repeating the experiment many times **probability theory i.e. statistics** is used to **estimate** what might have happened if the experiment had been repeated
- A mathematical model is used to fill this “data gap”
- Generally the most difficult task in statistics is to decide what “model” is most appropriate for a given experiment

# What is Statistics? - A definition

A set of analytical tools designed to quantify uncertainty

- If an experiment or procedure is repeated, how likely is it that the new results will be similar to those already observed?
- What is the likely variation in results if the experiment was repeated?



# What is Statistics? - A definition

The key scientific purpose of statistics

- to provide **evidence** of the existence of some “effect” of scientific interest
- i.e. evidence based medicine

## As a reminder: The importance of study design

Even the most sophisticated statistical analyses cannot rescue a poorly designed study

→ unreliable results

→ inability to answer the main research question

## Putting Statistics in Context

- The vast majority of analyses can be done in a straightforward fashion - just remember and always use common sense as a guide - be skeptical!
- It is very easy to get “lost” in the statistical software and technical jargon, which differs markedly between different software packages. Terminology can also differ greatly between textbooks...
- Wikipedia is as good a resource as any for finding out about different statistical tests and terminology

# Exploratory Data Analysis

- It is crucially important to explore your data fully before considering any “formal” statistical analyses

## Exploratory Data Analysis

- It is crucially important to explore your data fully before considering any “formal” statistical analyses
- What explorations are done depends on the objective of the study - the research question(s)

## Exploratory Data Analysis

- It is crucially important to explore your data fully before considering any “formal” statistical analyses
- What explorations are done depends on the objective of the study - the research question(s)
- Helps to decide what kind of formal statistical analyses might be most appropriate for the data available

# Exploratory Data Analysis

- It is crucially important to explore your data fully before considering any “formal” statistical analyses
- What explorations are done depends on the objective of the study - the research question(s)
- Helps to decide what kind of formal statistical analyses might be most appropriate for the data available
- What a simple descriptive analysis **does not** provide is evidence of whether the observed treatment effects are large enough to be notable once sampling variation has been accounted - that is the role of formal analyses, e.g. hypothesis testing

# Summary Statistics

## Continuous (Integers / Numeric)

- Mean - a measure of location. Always examine the average value of the response variable(s) for the different “treatment” effects in your data
- Median - a robust single value summary of a set of data (50% quantile point) - most useful in highly skewed data or data with outliers
- Standard deviation (sd) - a measure of spread, how variable the data are
- Standard error of the mean (se) - an estimate of how far the sample mean is likely to be from the population mean
- and others: min, max, range, IQR, ...





```
mean(x) # mean
```

```
median(x) # median
```

```
sd(x) # standard deviation
```

```
min(x) # minimum
```

```
max(x) # maximum
```

```
range(x) # range
```

```
IQR(x) # interquartile range
```

## Continuous Data Summaries

**standard deviation**

$$s = \sqrt{\frac{1}{n-1} \sum_{i=1}^n (x_i - \bar{x})^2}$$

**standard error**

$$se = \frac{s}{\sqrt{n}}$$

## Correlation coefficient

### Combination of continuous and continuous

Correlation coefficient a measure association between two continuous variables (common but somewhat limited)

#### Pearson's correlation coefficient r

$$r = \frac{\sum_{i=1}^n (X_i - \bar{X})(Y_i - \bar{Y})}{\sqrt{\sum_{i=1}^n (X_i - \bar{X})^2} \sqrt{\sum_{i=1}^n (Y_i - \bar{Y})^2}}$$

$\bar{X}$ : mean of variable x

$\bar{Y}$ : mean of variable y



```
# Test for Association/Correlation Between  
# Paired Samples  
cor.test(data$x, data$y, method = "pearson")  
cor.test(data$x, data$y, method = "spearman")  
  
# Scatterplot(s)  
pairs(data$x ~ data$y)  
pairs(data)
```

# Summary Statistics

## Continuous and ordinal variables



```
tapply(data$x.cont, data$y.fac, mean)
```

```
tapply(data$x.cont, data$y.fac, median)
```

```
tapply(data$x.cont, data$y.fac, sd)
```



- Median - a robust single value summary of a set of data (50% quantile point) - most useful in highly skewed data or data with outliers
- e.g. 10th and 90th percentile - a measure of spread, how variable the data are

```
quantile(x, probs = c(0.1, 0.9))
```

- proportions - e.g. percentage per grade

```
prop.table(table(data$x.fac))  
prop.table(table(data$x.fac, data$y.fac))
```



- proportions - percentage within the different categories
- contingency tables e. g.  $2 \times 2$

```
table(data$x.fac)
table(data$x.fac, data$y.fac)
prop.table(table(data$x.fac))
prop.table(table(data$x.fac, data$y.fac))
```

## Exercise 8





## How to deal with missing values in R? (1/3)

- In R, missing values are represented by the symbol **NA** (not available).
- Impossible values (e. g., dividing by zero) are represented by the symbol **NaN** (not a number).
- Ask yourself why a **NA** and / or **NaN** occurs!

# How to deal with missing values in R? (2/3)

- Testing for Missing Values

```
vec1 <- c(1, 2, 3, NA)
is.na(vec1) # returns a vector (FALSE, FALSE, FALSE, TRUE)
# The TRUE indicates the position of the NA in vec1.
```

- Recoding Values to Missing

```
# recode specific values (e. g. 0.001) to missing for variable x
# select rows where x is 0.001 and recode value in column x with NA
dat$x[dat$x == 0.001] <- NA
```

## How to deal with missing values in R? (3/3)

- Excluding Missing Values from specific function calls

```
a <- c(1, 2, NA, 3)
mean(a) # returns NA
mean(a, na.rm=TRUE) # returns 2
```

- Check for complete cases with function `complete.cases(...)`

```
# list rows of data that have missing values
dat[!complete.cases(dat),]
subdat <- dat[complete.cases(dat),]
```

- Create new dataset without missing data with function `na.omit(...)`

```
new.dat <- na.omit(dat)
```

# How to check your data for plausibility?

- Ask yourself what can go wrong?
- Implausible values?
- Impossible values?
- Logical errors?

## Exercise 9A: Plausibility Checks



## Exercise 9B: Missing Values



## Exercise 10



## Basic Statistical Tests

Study data is collected for a purpose - to answer one or more specific scientific questions. The classical way to perform a formal statistical analyses of these data is to formulate these research questions into statistical **hypothesis tests**.

In this section we will go through a simple example in detail to highlight some of the important concepts - the general approach for more complex analyses is exactly same. *Note: the precise technical details are much less important than the concepts!*



## Simple Example - One Population

**After six weeks will the mean weight of a chicken be more than 250 grams?**

There are 71 observations in `chickwts` from which to answer this question. This can be formulated into a statistical hypothesis test. A hypothesis test has two parts, the null hypothesis and the alternative hypothesis. This is typically written as follows:

$$H_0 : \mu \leq 250$$

$$H_A : \mu > 250$$

where  $\mu$  is the mean weight in the **population** of chickens from which the sample of 71 chickens was drawn. Remember - we know the mean weight in the sample of chickens is greater than 250 it is the **population** of chickens which we are interested in.

## Simple Example - One Population

**After six weeks will the mean weight of a chicken be at least 250 grams?**

$$H_0 : \mu \leq 250$$

$$H_A : \mu > 250$$

The null hypothesis ( $H_0$ ) is the default situation, sometimes called the “state of nature”. In a treatment-control trial,  $H_0$  is typically that the effect of the treatment is not different from the control. In this example our default position is that the mean weight of chickens is  $\leq 250$ . This is called a single-sided hypothesis test.



We now analyse the 71 observations to see whether there is evidence to **REJECT** the null hypothesis  $H_0$ , and if the null hypothesis is rejected then we can conclude that the available evidence supports the alternative hypothesis.

```
t.test(chickwts$weight, mu = 250)
```

Note that hypothesis testing is concerned with finding evidence in support of the null hypothesis  $H_0$  - the default situation - rather than evidence in favour of the alternative hypothesis.

## One Sample t-test

For the chicken weights data an appropriate formal analyses is to use a **one-sample t-test**, why this test is appropriate will be discussed later. This analysis involves calculating a simple summary statistic - called a *t*-statistic - which we do entirely from the observed data.

$$T_{obs} = \frac{\bar{x} - \mu}{s/\sqrt{n}}$$

where  $\bar{x}$  is the sample mean,  $s$  the sample standard deviation and  $\mu$  is the population mean in the null hypothesis which we wish to test for. We then look up the value of  $T_{obs}$  in a set of statistical tables/computer to see what the answer is to our research question.

## Important concept - sampling


Why is  $T_{obs} = \frac{\bar{x} - \mu}{s / \sqrt{(n)}}$  called a  $t$ -statistic?

If another sample of 71 chickens from the same population were weighed then the values for  $\bar{x}$  and  $s$  would be different, and hence the value for  $T_{obs}$ . If this was repeated many times and a histogram/Q-Q/P-P plot produced of the values for  $T_{obs}$  then this would follow the shape of a known distribution - **student-t probability distribution**. It is this piece of mathematics - knowing what the sampling distribution of  $T_{obs}$  is - which allows us to infer information about the population of chickens from which our original 71 chickens were sampled - without actually having to collect lots and lots of other samples of chickens! Mathematical theory is used to fill this data gap.

$$T_{obs} = \frac{261.31 - 250}{78.07 / \sqrt{71}} = 1.22$$

Put the values for the sample mean and standard deviation into the t-statistic formula along with the  $\mu = 250$ . We now look up the value of this in a t-distribution reference table. All this calculation will be done for you in R but it is important to understand the general process as this is the same for hypothesis testing in other more complex analyses.

# One Sample, one-sided, t-test



## Important concept - $p$ -values

- The result of a hypothesis test is usually communicated in the form of a  **$p$ -value**
- The interpretation of a  $p$ -value is of crucial importance - it is the *probability that the test statistic takes values at least as extreme as that observed **assuming that  $H_0$  is true***
- Exactly what **at least as extreme as** refers to depends on the alternative hypothesis  $H_A$ .
- This may sound rather abstract but it is usually obvious in practice



## Simple Example - One Population

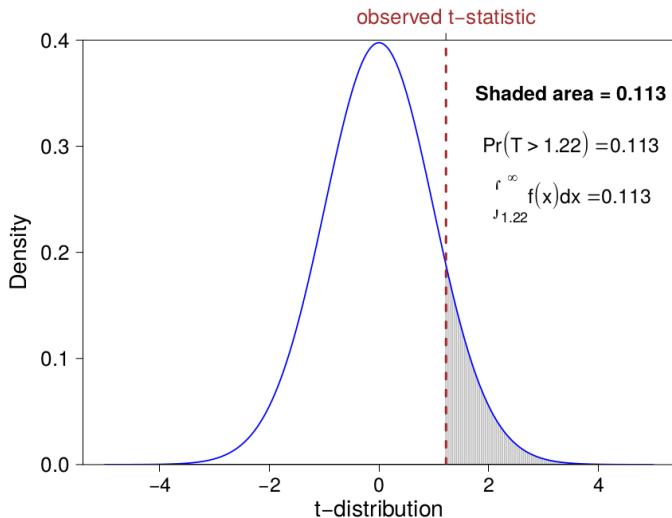
**After six weeks will the mean weight of a chicken be at least 250 grams?**

$$H_0 : \mu \leq 250$$

$$H_A : \mu > 250$$

The alternative hypothesis is  $\mu > 250$  so in this test **at least as extreme as** in the definition of a  $p$ -value is the probability of observing a t-statistic which is  $> 1.22$  **assuming that  $H_0$  is true** - this is why 250 was used for  $\mu$  when calculating  $T_{obs}$ .

# One Sample, one-sided, t-test



## Research Question

The purpose of this hypothesis test analysis is to answer a very specific scientific question:

**After six weeks will the mean weight of a chicken be more than 250 grams?**

So what is our answer?

The  $p$ -value for this hypothesis test is 0.113. Based on this value we can either **reject**  $H_0$  and conclude that the mean weight of chickens in the population is likely to be greater than 250 grams or else we can **accept**  $H_0$  where the mean chicken weight is less than 250 grams.

## Research Question - be pragmatic with $p$ -values

By convention a  $p$ -value of less than 0.05 is considered to provide reasonable evidence for rejecting  $H_0$ . A  $p$ -value of between 0.05 and 0.1 might be considered as weak evidence against  $H_0$ . Values of less than 0.01 are generally considered as very strong evidence for rejecting  $H_0$ . It is **always** best to provide a  $p$ -value in any analyses to let the reviewer/client see the strength of evidence rather than simply claiming statistically significant findings!

# Communicating Results of Hypothesis Tests

Transparency is essential - the devil can be in the detail - which at the very least should comprise:

- i. what hypothesis was being tested - be clear and precise
- ii. what statistical test was used
- iii. what the  $p$ -value is
- iv. what the treatment effect is (more later).

This is particularly crucial if the analyses are to be given to someone *e/se* to then make a judgment on the scientific significance.

## Two-sided Tests: One Population

**After six weeks will the mean weight of a chicken be equal to 250 grams?**

This is now a two sided hypothesis test:

$$H_0 : \mu = 250$$

$$H_A : \mu \neq 250$$

This time the hypothesis test is asking how much evidence is there in our sample data to conclude that in the population of all chickens the mean weight is not equal to 250 grams.

## Two-sided Tests: One Population



```
t.test(chickwts$weight, mu = 250)


##
##  One Sample t-test
##
## data:  chickwts$weight
## t = 1.2206, df = 70, p-value = 0.2263
## alternative hypothesis: true mean is not equal to 250
## 95 percent confidence interval:
##  242.8301 279.7896
## sample estimates:
## mean of x
##  261.3099
```

## Two-sided Tests

A two-sided test is similar to a one-sided test - the key difference is in what is now defined as **at least as extreme** in the definition of the  $p$ -value. This time the alternative hypothesis refers to observing a value of **either**  $\bar{x} > 250$  or  $\bar{x} < -250$  **assuming that  $H_0$  is true**, which using the t-test approach is equivalent to the probability of observing  $T_{obs} > 1.22$  or  $T_{obs} < -1.22$  which we can again look up in reference tables.



# One Sample, two-sided, t-test



## Two-sided Tests

- The two-sided t-test has a  $p$ -value which is exactly double the single sided test
- Think! - intuitively the  $p$ -value should be less for a single sided test as the research question you are asking is much narrower e.g. greater than 250 grams, as opposed to whether the mean chicken weight might be **either** less than 250 grams **or greater** than 250 grams.

→ You are using the same amount of information (71 observations) to answer a narrower research question and so all else being equal you should expect a “more powerful” analyses (e.g. a lower  $p$ -value all else being equal)

# Exercise 11



# Chisquare-Test

# The Chisquare Distribution

# Exact Fisher Test: Permutation Test

## Exercise 12



# Plotting in R







- Continuous data
  - Histogram
  - Boxplot
- Nominal / Ordinal data
  - Barplot
  - Mosaicplot
  - Scatterplots

## Exercise 13A and 13B



## Exercise 14



# Overview: ANOVA and linear models

- Introduction to ANOVA
- How to run an ANOVA in R
- Checking model assumptions in R
- Multiple comparisons options in R
- ANOVA as a special case of a linear model
- The simple linear regression model
- The multiple linear regression model
- Model selection:  $R^2$  and AIC
- Two-way Interactions in R (Day 6)
- Confounding (Day 6)

# Hypothesis Testing - One Way ANOVA

We have seen how to perform hypothesis tests when comparing two populations using the two sample t-test. In many analyses we may have multiple populations - for example suppose we have a treatment which has a number of different levels high/medium/low/placebo, or equivalently a number of different treatments. What then is the hypothesis we wish to test?

**Is there a difference in the effect of the treatment?**

$$H_0 : \mu_1 = \mu_2 = \cdots = \mu_k$$

$$H_A : \text{at least one pair of } \mu_1, \dots, \mu_k \text{ are different}$$

where  $\mu_1, \dots, \mu_k$  denote the mean effect of treatment levels 1 through  $k$ .

# One Way ANOVA

Analysis of variance, **ANOVA**, to analyze differences between group means. The observed variance in the outcome variable is partitioned into components attributable to different sources of variation.

\*\*\*\*\*

## ANOVA estimates three sample variances (sum of squares)

- a total variance based on all observation deviations from the grand mean
- a variance based on the group mean deviations from the grand mean
- an (error) variance based on all the observation deviations from their group mean




An **F-test**, a statistical test, in which the test statistic has an F-distribution under the null hypothesis is used to assess statistical significance in an ANOVA.

- degrees of freedom

## ANOVA in R with chickwts



```
data(chickwts)  
boxplot(chickwts$weight ~ chickwts$feed)
```



```
# aov.mod <- aov(chickwts$weight ~ chickwts$feed)
aov.mod <- aov(weight ~ feed, data = chickwts)
# What objects can we extract from a anova model?
objects(aov.mod)

## [1] "assign"          "call"            "coefficients"    "contrasts"
## [5] "df.residual"     "effects"         "fitted.values"   "model"
## [9] "qr"              "rank"            "residuals"       "terms"
## [13] "xlevels"

#
summary(aov.mod)

##              Df Sum Sq Mean Sq F value    Pr(>F)
## feed           5 231129   46226    15.37 5.94e-10 ***
## Residuals     65 195556    3009
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

## Overview: Checking model assumptions

- $\text{mean}(\text{residuals}) = 0$
- Residuals are normally distributed (qqnorm, qqplot)
- Variance homoscedasticity (Bartlett & Levene's Test)
- Cook's distance: Influential data points
- Any pattern(s)?

## Checking model residuals: $\text{mean}(\text{residuals}) = 0$



"Unexplained rest of the model"

```
# What are residuals?
chickwts$residuals <- residuals(aov.mod)
tapply(chickwts$weight, chickwts$feed, mean)

##      casein horsebean  linseed  meatmeal   soybean sunflower
## 323.5833  160.2000  218.7500  276.9091  246.4286  328.9167

chickwts[c(1:3),]

##      weight      feed residuals
## 1      179 horsebean      18.8
## 2      160 horsebean      -0.2
## 3      136 horsebean     -24.2

# Save residuals to an objects and check mean of residuals
aov.mod.resid <- residuals(aov.mod)
mean(aov.mod.resid)


## [1] 7.573045e-16
```

# Checking model residuals: Residuals normally distributed

"Unexplained rest of the model"



```
par(mfrow=c(1,1))
qqnorm(aov.mod.resid)
qqline(aov.mod.resid, col = "red", lwd = 3, lty = 2)
```



*# Shapiro-Wilk test (dependent on sample size --> limited use)*

```
shapiro.test(aov.mod.resid)
```

```
##
```

```
## Shapiro-Wilk normality test
```

```
##
```

```
## data: aov.mod.resid
```

```
## W = 0.98616, p-value = 0.6272
```

# Checking model residuals: Variance homoscedasticity (1/3)

## Hypothesis tests



```
# Bartlett Test
bartlett.test(chickwts$weight ~ chickwts$feed)

##
## Bartlett test of homogeneity of variances
##
## data:  chickwts$weight by chickwts$feed
## Bartlett's K-squared = 3.2597, df = 5, p-value = 0.66


# Levene's Test
# library("Rcmdr")
# levene.test(chickwts$weight ~ chickwts$feed)
```

## Checking model residuals: Variance homoscedasticity (2/3)

### Graphical interpretation is better!



```
# Plot fitted against residual values  
plot(fitted.values(aov.mod), residuals(aov.mod))
```






## Checking model residuals: Variance homoscedasticity (3/3)

### Graphical interpretation is better!




```
# Plot fitted against residual values
par(mfrow=c(1,2), pty="s", mar = c(10, 4, 1, 2))
plot(fitted.values(aov.mod), residuals(aov.mod))
abline(h = 0, col = "red", lwd = 3, lty = 2)
plot(aov.mod, which=1)
```



# Checking model residuals: Cook's distance






```
# Plot fitted against residual values  
# Cut-off at 3  
par(mfrow=c(1,1), pty="s", mar = c(5, 4, 4, 2))  
plot(aov.mod, which=4)
```



# Checking for potential patterns





```
par(mfrow=c(1,3), pty="s", mar = c(10, 4, 1, 2))
plot(fitted.values(aov.mod), residuals(aov.mod))
abline(h = 0, col = "red", lwd = 3, lty = 2)
# Plot residuals against variables from the model
plot(chickwts$weight, residuals(aov.mod), ylab = "residuals")
plot(chickwts$feed, residuals(aov.mod),
     xlab = "chickwts$feed", ylab = "residuals")
```



# Plot anova objects



```
par(mfrow=c(2, 2))  
plot(aov.mod)
```





So far, we know there is difference between the feed types.

However, we do not yet know which feed type differ.

In principal, we could do multiple t-tests, BUT ... we would use our data several times. Classically, we choose an  $\alpha$ -level of 5 %, in cases of multiple comparisons we are facing the familywise error rate:

$$FWE \leq 1 - (1 - \alpha)^n,$$

where  $\alpha = 5 \%$  and  $n =$  number of comparisons.

```
alpha = 0.05  
1 - (1 - alpha)^1    # n = 1    # [1] 0.05  
1 - (1 - alpha)^5    # n = 5    # [1] 0.2262191  
1 - (1 - alpha)^10   # n = 10   # [1] 0.4012631
```



Hence, we are better of using one of the following procedures to adjust for multiple comparisons:

- **Bonferroni:**  $p$ -value correction by testing each individual hypothesis at a significance level of  $\frac{\alpha}{m}$ , where  $\alpha$  is the desired overall  $\alpha$  level and  $m$  is the number of hypotheses.
- **Dunnett:** Multiple comparisons of each group to a reference.
- **Tukey Honest Significant Differences** (Homogeneous subgroups): Multiple comparisons of all possible combinations.
- ...

# Multiple comparisons options in R: Bonferroni



```
aov.mod <- aov(weight ~ feed, data = chickwts)
pairwise.t.test(chickwts$weight, chickwts$feed, p.adj = "none")

##
## Pairwise comparisons using t tests with pooled SD
##
## data:  chickwts$weight and chickwts$feed
##
##           casein  horsebean  linseed  meatmeal  soybean
## horsebean 2.1e-09 -          -          -          -
## linseed   1.5e-05 0.01522   -          -          -
## meatmeal  0.04557 7.5e-06   0.01348 -          -
## soybean   0.00067 0.00032   0.20414 0.17255  -
## sunflower 0.81249 8.2e-10   6.2e-06 0.02644 0.00030
##
## P value adjustment method: none

pairwise.t.test(chickwts$weight, chickwts$feed, p.adj = "bonferroni")

##
## Pairwise comparisons using t tests with pooled SD
##
## data:  chickwts$weight and chickwts$feed
##
##           casein  horsebean  linseed  meatmeal  soybean
## horsebean 3.1e-08 -          -          -          -
## linseed   0.00022 0.22833   -          -          -
## meatmeal  0.68350 0.00011   0.20218 -          -
## soybean   0.00998 0.00487   1.00000 1.00000  -
## sunflower 1.00000 1.2e-08   9.3e-05 0.39653 0.00447
##
## P value adjustment method: bonferroni
```

\*\*\*\*\*





# Multiple comparisons options in R: Dunnett



```
library("multcomp")
# compares always to baseline levels (here: casein) --> saves degrees of freedom
dunnett <- glht(aov.mod, linfct = mcp(feed = "Dunnett"))
summary(dunnett)

##
## Simultaneous Tests for General Linear Hypotheses
##
## Multiple Comparisons of Means: Dunnett Contrasts
##
##
## Fit: aov(formula = weight ~ feed, data = chickwts)
##
## Linear Hypotheses:
##
##           Estimate Std. Error t value Pr(>|t|)
## horsebean - casein == 0 -163.383    23.485  -6.957 < 0.001 ***
## linseed - casein == 0  -104.833    22.393  -4.682 < 0.001 ***
## meatmeal - casein == 0  -46.674    22.896  -2.039 0.16716
## soybean - casein == 0   -77.155    21.578  -3.576 0.00308 **
## sunflower - casein == 0    5.333    22.393   0.238 0.99945
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
## (Adjusted p values reported -- single-step method)
```

# Multiple comparisons options in R: Tukey

## Tukey Honest Significant Differences



```
library("multcomp")
# compares all factor levels
tukey <- glht(aov.mod, linfct = mcp(feed = "Tukey"))
summary(tukey)

##
## Simultaneous Tests for General Linear Hypotheses
##
## Multiple Comparisons of Means: Tukey Contrasts
##
##
## Fit: aov(formula = weight ~ feed, data = chickwts)
##
## Linear Hypotheses:
##
##      Estimate Std. Error t value Pr(>|t|)
## horsebean - casein == 0 -163.383    23.485  -6.957 < 0.001 ***
## linseed - casein == 0  -104.833    22.393  -4.682 < 0.001 ***
## meatmeal - casein == 0  -46.674    22.896  -2.039 0.03195
## soybean - casein == 0   -77.155    21.578  -3.576 0.00837 **
## sunflower - casein == 0    5.333    22.393    0.238 0.99989
## linseed - horsebean == 0  58.550    23.485    2.493 0.14100
## meatmeal - horsebean == 0 116.709    23.966    4.870 < 0.001 ***
## soybean - horsebean == 0  86.229    22.710    3.797 0.00425 **
## sunflower - horsebean == 0 168.717    23.485    7.184 < 0.001 ***
## meatmeal - linseed == 0   58.159    22.896    2.540 0.12741
## soybean - linseed == 0    27.679    21.578    1.283 0.79299
## sunflower - linseed == 0  110.167    22.393    4.920 < 0.001 ***
## soybean - meatmeal == 0   -30.481    22.100   -1.379 0.73881
## sunflower - meatmeal == 0  52.008    22.896    2.271 0.22032
## sunflower - soybean == 0   82.488    21.578    3.823 0.00385 **
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```


Scanned with CamScanner reported -- single-step method) Lecture Slides (all)

# Multiple comparisons options in R: Tukey

## Tukey Honest Significant Differences with homogeneous subgroups



```
# summary(tukey)           # standard display
tukey.cld <- cld(tukey)    # letter-based display
# the cld(...) function sets up a compact letter display of all pair-wise comparisons
par(mfrow=c(1,1), mar=c(5, 4, 5, 2))
plot(tukey.cld)
```



## Exercise 15



# Introduction: ANOVA as a special case of a linear model

## Linear Modelling

So far we have considered how to determine whether statistically significant differences exist between different "feed" groups (factors).

The explanatory variable has been categorical. We have not yet considered continuous explanatory variables.

ANOVA and linear regression are both a special cases of **linear models**.

## Simple linear regression (1/3)

A **simple linear regression** fits a straight line through a set of data points. This straight line is fitted in a way that makes the sum of the squared residuals (the vertical distances between each data point and the fitted line) as small as possible.

## Simple linear regression (2/3)

In a simple linear regression model for  $n$  data points  $(x_i, y_i)$ ,  $i = 1, \dots, n$  the following equation is used:

$$y_i = \alpha + \beta x_i + \epsilon,$$

where

- $\alpha$ : intercept or constant
- $\beta$ : slope, regression coefficient (effect size)
- $\epsilon$ : error, residuals

## Simple linear regression (3/3)

The goal is to find the equation of the straight line

$$f(x) = y = \alpha + \beta x,$$

which would provide a "best" fit for the data points (least square approach).



# Assumptions of a linear regression

A **simple linear regression** is based on several assumptions which should be checked carefully.

- linearity of the relationship between the explanatory (independent) and the outcome (dependent) variable
- normality of the residuals
- independence
- constant variance (homoscedasticity)



```
## Error in library("HSAUR3"): there is no package called 'HSAUR3'  
## Error in str(water): object 'water' not found
```

```
data(water)
```

```
## Warning in data(water): data set 'water' not found
```

```
# mod.hard <- lm(mortality ~ hardness, data = water)  
mod.hard <- lm(water$mortality ~ water$hardness)
```

```
## Error in eval(predvars, data, env): object 'water' not found
```

```
summary(mod.hard)
```

```
## Error in summary(mod.hard): object 'mod.hard' not found
```

# Checking linear model assumptions



```
par(mfrow=c(2,2))  
plot(mod.hard)
```

```
## Error in plot(mod.hard): object 'mod.hard' not found
```

# Linear model vs. t.test(...) in R



```
mod.loc <- lm(mortality ~ location, data = water)

## Error in is.data.frame(data):  object 'water' not found

coef(mod.loc)

## Error in coef(mod.loc):  object 'mod.loc' not found

t.test(water$mortality ~ water$location)

## Error in eval(predvars, data, env):  object 'water' not found
```

# The multiple linear model

## An extension of the simple linear model (1/2)

If several explanatory variables are of interest, instead of performing multiple simple regressions, a **multiple linear regression** or **multivariable** approach is appropriate.

- the same assumptions as for simple linear regressions should be checked
- collinearity might be an issue (see `vif(...)` function from package `usdm`)
- model comparison (AIC) (...discussed later)

# The multiple linear model

## An extension of the simple linear model (2/2)

Simple linear regression model:

$$y_i = \alpha + \beta x_i + \epsilon$$

Multiple linear regression model:

$$y_i = \alpha + \beta_1 x_{1i} + \beta_2 x_{2i} + \dots + \epsilon, \quad (1)$$

$$y_i = \alpha + \beta_1 x_{1i} * \beta_2 x_{2i} + \dots + \epsilon,$$

where

- $\alpha$ : intercept or constant
- $\beta_1$ : slope1, regression coefficient (effect size)
- $\beta_2$ : slope2, regression coefficient (effect size)
- $\beta_1 * \beta_2$ : interaction effect between  $x_1$  and  $x_2$
- $\epsilon$ : error, residuals

# The multiple linear model

## Interpretation of model coefficients

- $\alpha$ : intercept or constant
  - $\beta_1$ : slope1, regression coefficient (effect size)
  - $\beta_2$ : slope2, regression coefficient (effect size)
  - $\beta_1 * \beta_2$ : interaction effect between  $x_1$  and  $x_2$
  - $\epsilon$ : error, residuals
- 
- $\rightarrow \beta_1$  describes the number of units of a change in the outcome variable  $y$  as  $x_1$  changes by one unit,  $x_2$  being held constant.
  - $\rightarrow \beta_2$  describes the number of units of a change in the outcome variable  $y$  as  $x_2$  changes by one unit,  $x_1$  being held constant.



```
mod.hard.loc <- lm(mortality ~ hardness + location, data = water)

## Error in is.data.frame(data):  object 'water' not found

summary(mod.hard.loc)

## Error in summary(mod.hard.loc):  object 'mod.hard.loc' not found
```



\*\*\*\*\*





- Multiple  $R^2$ : the larger, the better
- Akaike criterion (AIC): the smaller, the better



Two models are nested if one of them is a particular case of the other one: the simpler model can be obtained by setting some coefficients of the more complex model to particular values.

```
mod.hard <- lm(mortality ~ hardness, data = water)           # mod 1

## Error in is.data.frame(data):  object 'water' not found

mod.loc <- lm(mortality ~ location, data = water)            # mod 2

## Error in is.data.frame(data):  object 'water' not found

mod.hard.loc <- lm(mortality ~ hardness + location, data = water) # mod 3

## Error in is.data.frame(data):  object 'water' not found
```

Among the 3 above models...

- which ones are nested?
- which ones are not nested?



$R^2$  is a measure of fit quality:

$$R^2 = \frac{\text{SSRegression}}{\text{SSTotal}} \quad (2)$$
$$R^2 = 1 - \frac{\text{SSError}}{\text{SSTotal}}$$

## CAREFUL:

- The SSError always decreases as more predictors are added to the model.
- $R^2$  always increases and can be artificially large.

## Akaike criterion (AIC)

- Model fit ( $R^2$ ) always improves with model complexity. We would like to strike a good balance between **model fit** and **model simplicity**.
- AIC combines a measure of model fit with a measure of model complexity: The smaller, the better.

For a given data set and a given model

$$AIC = -2\log(L) + 2p,$$

where  $-2\log(L)$  is a function of the prediction error: the smaller, the better. Measures how the model fits the data.  $2p$  penalizes complex models: the smaller, the better.

## Model selection strategy for AIC

- Consider a number of candidate models.  
(They need not be nested.)
- Calculate their AIC.
- Choose the model(s) with the smallest AIC.

→ **CAREFUL**: The absolute value of AIC is meaningless. The relative AIC values, between models, is meaningful.



```
mod1 <- lm(mortality ~ hardness, data = water)

## Error in is.data.frame(data):  object 'water' not found

mod2 <- lm(mortality ~ location, data = water)

## Error in is.data.frame(data):  object 'water' not found

mod3 <- lm(mortality ~ hardness + location, data = water)

## Error in is.data.frame(data):  object 'water' not found

AIC(mod1, mod2, mod3)

## Error in AIC(mod1, mod2, mod3):  object 'mod1' not found

round(AIC(mod1, mod2, mod3), 2)

## Error in AIC(mod1, mod2, mod3):  object 'mod1' not found
```



```
mod3 <- lm(mortality ~ hardness + location, data = water)

## Error in is.data.frame(data): object 'water' not found

par(mfrow=c(2,2))
plot(mod3)

## Error in plot(mod3): object 'mod3' not found
```



## Exercise 16




## Exercise 17



# Interactions



```
interaction.plot(ToothGrowth$dose, ToothGrowth$supp, ToothGrowth$len,  
                fixed = TRUE)
```





There are three different interactions:

- **Interaction between two categorical variables**
- Interaction between one continuous and one categorical variables
- Interaction between two continuous variables



Model specification in R: Be aware, an interaction is never tested without its corresponding main effects included in the model.

```
# Interaction between two categorical variables
# mod.dose.supp.int <- lm(len ~ dose.fac + supp + dose.fac:supp,
# data = ToothGrowth)
mod.dose.supp.int <- lm(len ~ dose.fac * supp, data = ToothGrowth)
# summary(mod.dose.supp.int)
```

$$\begin{aligned} y \sim & \beta_{\text{baseline}((\text{dose}==\text{low}) \& (\text{supp}==\text{OJ}))} + \beta_{\text{dose}==\text{med}} + \\ & \beta_{\text{dose}==\text{high}} + \beta_{\text{supp}==\text{VC}} \\ & + \beta_{(\text{dose}==\text{med}) \& (\text{supp}==\text{VC})} \\ & + \beta_{(\text{dose}==\text{high}) \& (\text{supp}==\text{VC})} \end{aligned}$$

## Two-way Interactions in R (3/3)



```
# Interaction between two categorical variables
# mod.dose.supp.int <- lm(len ~ dose.fac + supp + dose.fac:supp,
# data = ToothGrowth)
mod.dose.supp.int <- lm(len ~ dose.fac * supp, data = ToothGrowth)
summary(mod.dose.supp.int)

##
## Call:
## lm(formula = len ~ dose.fac * supp, data = ToothGrowth)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -8.20  -2.72  -0.27   2.65   8.27
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)      13.230      1.148   11.521 3.60e-16 ***
## dose.faced       9.470       1.624    5.831 3.18e-07 ***
## dose.fachigh     12.830       1.624    7.900 1.43e-10 ***
## suppVC          -5.250       1.624   -3.233 0.00209 **
## dose.faced:suppVC -0.680       2.297   -0.296 0.76831
## dose.fachigh:suppVC 5.330       2.297    2.321 0.02411 *
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 3.631 on 54 degrees of freedom
## Multiple R-squared:  0.7937, Adjusted R-squared:  0.7746
## F-statistic: 41.56 on 5 and 54 DF, p-value: < 2.2e-16
```

## Two-way Interactions in R: Interpretation of coefficients (1/2)



```
coef(mod.dose.supp.int)
```

```
##          (Intercept)          dose.facmed          dose.fachigh  
##           13.23           9.47           12.83  
##           suppVC  dose.facmed:suppVC  dose.fachigh:suppVC  
##           -5.25           -0.68           5.33
```

$$\begin{aligned} y \sim & \beta_{\text{baseline}((\text{dose}==\text{low}) \& (\text{supp}==\text{OJ}))} + \beta_{\text{dose}==\text{med}} + \\ & \beta_{\text{dose}==\text{high}} + \beta_{\text{supp}==\text{VC}} \\ & + \beta_{(\text{dose}==\text{med}) \& (\text{supp}==\text{VC})} \\ & + \beta_{(\text{dose}==\text{high}) \& (\text{supp}==\text{VC})} \end{aligned}$$



```
coef(mod.dose.supp.int)
```

```
##          (Intercept)          dose.facmed          dose.fachigh  
##          13.23          9.47          12.83  
##          suppVC  dose.facmed:suppVC  dose.fachigh:suppVC  
##          -5.25          -0.68          5.33
```

- → `dose.facmed:suppVC`  
change in the slope between the low and med dose.fac group under the supplement type (supp) VC in comparison to the intercept. In other words, by changing the dose from low to med within the supplement group VC, the slope **decreases** by approx.  $-0.68$ .
- → `dose.fachigh:suppVC`  
change in the slope between the low and high dose.fac group under the supplement type (supp) VC in comparison to the intercept. In other words, by changing the dose from low to high within the supplement group VC, the slope **increases** by approx.  $+5.33$ .



\*\*\*\*\*



## Exercise 18



- **Confounding variable** - is associated with the exposure and the outcome
- **Confounding variable** - is not part of the causal path between exposure and the outcome

confounding.png

## Some things to remember

- Not adjusting / controlling for a confounding variable may lead to bias results. It is good practice to present as well the crude (not adjusted ORs) and the adjusted ones. Adjustment is typically done if difference  $> 10\%$ .
- Check also for interaction (= effect modification) e. g. using logistic regression.
- Do not adjust for a variable C if it is a common effect of E and D (collider) or if it is in the causal pathway of E and D.

## Other topics

So far, we have covered some of the basic **tools** necessary for helping you to start to analyse your own study data.

Many things have been mentioned others have not been included - some are worth mentioning as they are particularly common in some areas of veterinary research.

## Normally distributed response variables

So far, we have only considered analyses where the response variable is a **continuous** measurement.

We now have a brief look at two other common types of **response** variables.

The methods we use are largely similar to what we have looked at previously.



```
## Error in file(file, "rt"): cannot open the connection
## Error in '$<-.data.frame'('*tmp*', respsymptoms, value =
structure(integer(0), .Label = c("no", : replacement has 0 rows, data
has 228
```



- numeric & integers

```
ToothGrowth$len[1:6]
## [1]  4.2 11.5  7.3  5.8  6.4 10.0

bacteria$week[1:6]
## [1]  0  2  4 11  0  2
```





- numeric & integers

```
ToothGrowth$len[1:6]
## [1]  4.2 11.5  7.3  5.8  6.4 10.0

bacteria$week[1:6]
## [1]  0  2  4 11  0  2
```

- unordered factor with 2 levels

```
lung$sex[1:6]
## [1] male male male male male male
## Levels: female male
```

- (un/ordered) factor (with more than 2 levels)

```
chickwts$feed[1:6]
## [1] horsebean horsebean horsebean horsebean horsebean horsebean
## Levels: casein horsebean linseed meatmeal soybean sunflower
```

# Generalised Linear Modelling

- random component:  $\mathbf{E}(\mathbf{Y}) = \mu$
- systematic component: covariates:  $\mathbf{x}_1, \mathbf{x}_2, \dots, \mathbf{x}_p$  produce a linear predictor  $\eta = \sum_1^p \mathbf{x}_j \beta_j$
- link between random and systematic components:  $\mu = \eta$ , e.g.
  - **Logistic Regression:**  
 $\eta_i = g(\mu_i)$ ,  $g(\cdot)$  is logit function for binomial data
  - **Poisson Regression:**  
 $\eta_i = g(\mu_i)$ ,  $g(\cdot)$  is exponential function for poisson data



E. g. a logistic regression with one covariate (`sex`) fitting a model to  $\mathbb{P}(Y = \text{TRUE}) = \pi$ , e. g. presence of respiratory symptoms (experienced by the child over the previous 12 months) , gives

$$\log \left( \frac{\pi}{1 - \pi} \right) = \beta_0 + \beta_1 \mathbf{x}_1$$

```
logreg.mod <- glm(respsymptoms ~ sex,  
                  family = binomial, # description of the link function  
                  data = lung)  
summary(logreg.mod)
```

# Logistic Regression: unordered factor with 2 levels as response



```
logreg.mod <- glm(respsymptoms ~ sex,  
                 family = binomial, # description of the link function  
                 data = lung)
```

```
## Error in eval(predvars, data, env): object 'respsymptoms' not found
```

```
summary(logreg.mod)
```

```
## Error in summary(logreg.mod): object 'logreg.mod' not found
```

# Logistic Regression: unordered factor with 2 levels as response

## Interpretation of model coefficients



```
summary(logreg.mod)


## Error in summary(logreg.mod): object 'logreg.mod' not found

cbind(logodds = coef(logreg.mod), odds = exp(coef(logreg.mod)))

## Error in coef(logreg.mod): object 'logreg.mod' not found
```

## Odds ratios versus log odds ratios

- The logistic regression model in R outputs the log odds ratios.
- $\exp(\text{log odds ratio}) = \text{odds ratio}$
- An **odds ratio (OR)** is a measure of association between an exposure (here: `sex`) and an outcome (here: `respsymptoms`). The OR represents the odds that an outcome will occur given a particular exposure, compared to the odds of the outcome occurring in the absence of that exposure.



# Logistic Regression: unordered factor with 2 levels as response

## Interpretation of model coefficients (cont'd)



```
cbind(logodds = coef(logreg.mod), odds = exp(coef(logreg.mod)))  
  
## Error in coef(logreg.mod): object 'logreg.mod' not found
```

The coefficient  $\beta_{\text{sex}}$  is the estimated amount by which the log odds of respsymptoms would in-/decrease if sex is male. The log odds of respsymptoms when sex is female is just above in the first row ( $\beta_{(\text{Intercept})}$ , reference level of sex).

## Logistic Regression: probability of success as response



- A researcher is examining beetle mortality after 5 hours of exposure to carbon disulphide, at various levels of concentration of the gas.
- Beetles were exposed to gaseous carbon disulphide at various concentrations (in mg/L) for five hours (Bliss, 1935) and the number of beetles killed were noted.

```
library("investr")

## Error in library("investr"):  there is no package called 'investr'

data(beetle)

## Warning in data(beetle):  data set 'beetle' not found

colnames(beetle) <- c("Dose", "Num.Beetles", "Num.Killed")

## Error in colnames(beetle) <- c("Dose", "Num.Beetles", "Num.Killed"):
object 'beetle' not found

str(beetle)
```



# Logistic Regression: probability of success as response



```
logreg.mod <- glm(cbind(Num.Killed, Num.Surv) ~ Dose,  
                  family = binomial, # description of the link function  
                  data = beetle)
```

```
## Error in is.data.frame(data): object 'beetle' not found
```

```
summary(logreg.mod)
```

```
## Error in summary(logreg.mod): object 'logreg.mod' not found
```

# Logistic Regression: probability of success as response

## Interpretation of model coefficients



```
summary(logreg.mod)

## Error in summary(logreg.mod): object 'logreg.mod' not found

cbind(logodds = coef(logreg.mod), odds = exp(coef(logreg.mod)))

## Error in coef(logreg.mod): object 'logreg.mod' not found
```

# Logistic Regression: probability of success as response

## Interpretation of model coefficients (cont'd)



```
cbind(logodds = coef(logreg.mod), odds = exp(coef(logreg.mod)))  
  
## Error in coef(logreg.mod): object 'logreg.mod' not found
```

The coefficient  $\beta_{\text{Dose}}$  is the estimated amount by which the log odds of being killed would in-/decrease if Dose would increase by one unit. The log odds of being killed when Dose is 0 is just above in the first row ( $\beta_{(\text{Intercept})}$ ).

## Some things to remember

- Using generalised linear models (glm) is very similar to the previous models we have seen for normal data (lm) - things like residuals need to be checked for randomness, although this is more difficult with counts. Also the residuals being normally distributed is not relevant to these models.
- **As always explore the data first visually and with tables before doing any formal analyses.**

## Short excursion to quantile regression

- When talking about regression, we almost exclusively model the conditional mean of a response given one or more explanatory variables.
- With the classical linear model we cannot skewed or otherwise non normal distribution as the corresponding quantiles from the linear model will be misleading. Therefore, we shift our attention to a completely distribution free approach that directly addresses conditional quantile modelling.  
(Text taken from Hothorn, Torsten, and Brian S. Everitt. A handbook of statistical analyses using R. CRC press, 2014.)

## Short excursion to quantile regression

The data set `sambia.csv` (given on the switch drive) contains information on malnutrition of sambian children. The outcome variable is the *z-score*, which gives information on the severity of malnutrition.

The following covariates are considered:

---

<code>zscore</code>	Z-score of the child
<code>sex</code>	Sex of the child (0 = female, 1 = male)
<code>age.child</code>	Age of the child
<code>work</code>	Is the mother working (0 = no, 1 = yes)
<code>age.mother.birth</code>	Age of the mother at birth of child
<code>bmi</code>	BMI of the mother

---

A *z-score* indicates how many standard deviations an element is from the mean. A *z-score* can be calculated from the following formula.  $z = (X - \mu) / \sigma$  where  $z$  is the *z-score*,  $X$  is the value of the element,  $\mu$  is the population mean, and  $\sigma$  is the standard deviation.



```
## Warning in file(file, "rt"): cannot open file
' /home/mburi/Dropbox/201710_Makerere/03_Exercises/data/sambia.csv': No
such file or directory
## Error in file(file, "rt"): cannot open the connection
## Error in factor(sambia$sex, levels = c(0, 1), labels = c("female",
"male")): object 'sambia' not found
## Error in factor(sambia$work, levels = c(0, 1), labels = c("no",
"yes")): object 'sambia' not found
```



```
sambia <- read.csv("~/201710_Makerere/data/sambia.csv", sep = ",")
str(sambia)
sambia$sex <- factor(sambia$sex, levels = c(0,1),
                    labels = c("female", "male"))
sambia$work <- factor(sambia$work, levels = c(0,1),
                    labels = c("no", "yes"))
```

```
str(sambia)
```

```
## Error in str(sambia): object 'sambia' not found
```



## Quantile regression in R (cont'd)



```
library("quantreg")

## Loading required package: SparseM
##
## Attaching package: 'SparseM'
## The following object is masked from 'package:base':
##
##   backsolve
##
## Attaching package: 'quantreg'
## The following object is masked from 'package:survival':
##
##   untangle.specials

# tau = 2.5 %
rq.025 <- rq(zscore ~ sex + age.child + work + bmi + age.mother.birth,
             tau = 0.025,
             data = sambia)

## Error in is.data.frame(data): object 'sambia' not found

# tau = 5 %
rq.05 <- rq(zscore ~ sex + age.child + work + bmi + age.mother.birth,
            tau = 0.05,
            data = sambia)

## Error in is.data.frame(data): object 'sambia' not found

# tau = 50 %
rq.5 <- rq(zscore ~ sex + age.child + work + bmi + age.mother.birth,
           tau = 0.5,
           data = sambia)
```

## Quantile regression in R (cont'd)

### Comparison of the three models `rq.025`, `rq.05`, `rq.5`



```
rq_coef <- cbind(coef(rq.025), coef(rq.05), coef(rq.5))

## Error in coef(rq.025): object 'rq.025' not found

rownames(rq_coef) <- rownames(summary(rq.5)$coefficients)

## Error in summary(rq.5): object 'rq.5' not found

colnames(rq_coef) <- paste0("tau = ", c("0.025", "0.05", "0.5"))

## Error in colnames(rq_coef) <- paste0("tau = ", c("0.025", "0.05", "0.5")): object 'rq_coef'
not found

round(rq_coef, 3)

## Error in eval(expr, envir, enclos): object 'rq_coef' not found
```

# Quantile regression in R (cont'd)

## Estimate $r_q$ for a grid of quantiles



```
# Estimate rq for a single tau
# rq.5 <- rq(zscore ~ sex + age.child + work + bmi + age.mother.birth,
#           tau = 0.5,
#           data = sambia)
# Estimate rq for a grid of tau
rq <- rq(zscore ~ sex + age.child + work + bmi + age.mother.birth,
        tau = seq(from = 0.025, to = 0.5, by = 0.025),
        data = sambia)

## Error in is.data.frame(data): object 'sambia' not found
```

# Quantile regression in R (cont'd)

## Plot results



```
plot(rq) # plot including line for the OLS coefficient (as estimated by lm)
```

```
## Error in formula.default(object, env = baseenv()): invalid formula
```

# Quantile regression in R (cont'd)

## Plot results in more details



```
# A sequence of coefficient estimates for quantile regressions with  
# varying tau parameters is visualized along with associated confidence bands.  
plot(summary(rq))
```

```
## Error in object[[i]]: object of type 'closure' is not subsettable
```

## Acknowledgement

We would like to thank **Beate Sick** for the inspiration given for these slides on *missing values* and *multiple imputation*.

- **Nonresponse** occurs when information on all variables is missing for a subset of the sample. Such non-responders may cause **selection bias**, as they effectively modify the study population to include only respondents.
- **Item nonresponse** occurs when only some variables have missing responses for a given individual.
  - If missing probabilities of a predictor do not depend on the outcome, then X is missing **nondifferentially**. Association between outcome and missingness of a predictor can be tested!
  - **Differential missingness probabilities** depend on disease outcome. This mainly arise in retrospective (e. g. case-control) studies. If predictors (exposure) are measured prospectively before outcome such dependency cannot occur in a direct way.

## Example: A study of HIV-positive people in NYC



The CHAIN project was a longitudinal cohort study of people living with HIV in New York City, which was recruited in 1994. Here we study a data subset collected from 532 subjects in the sixth round of interviews.

- `log_virus`: Log of self reported viral load level (=0 if below detection limit).
- `age`: The respondent's age at time of interview.
- `income`: The respondent's family annual income.
- `healthy`: A continuous scale of physical health (0-100).
- `mental`: Measure of poor mental health: 0/1 = No/Yes.
- `damage`: Ordered interval (0 to 5) for CD4 count.
- `treatment`: A three-level-ordered variable (0 = not currently taking HAART therapy, 1 = taking HAART nonadherent, 2 = taking HAART adherent)





```
## Error in library("mi"):  there is no package  
called 'mi'  
## Error in eval(expr, envir, enclos):  object  
'CHAIN' not found  
## Error in factor(chain$mental, levels = c(0, 1),  
labels = c("no", "yes")):  object 'chain' not found  
## Error in factor(chain$treatment, levels = c(0, 1,  
2), labels = c("0", :  object 'chain' not found
```

```
install.packages("mi")  
library("mi")  
data("CHAIN")  
chain <- CHAIN # rename the data frame  
names(chain)  
str(chain)  
summary(chain)
```

# Descriptive Statistics of the CHAIN data set in R




```
hist(chain$log_virus, ylim = c(0, 200), col = "darkgray")
nrow(chain) # [1] 532
apply(chain, 2, function(x) {sum(is.na(x))})
# log_virus      age      income    healthy    mental    damage treatment
#      179        24        38        24        24        63        24
summary(chain)
```

## Reasons for Item Nonresponse

- Individuals may not respond to specific questions on sensitive issues (e. g. questions on household income, drug abuse, ...).
- Medical or occupational records may be incomplete. Material might be spilled or reports might be lost in the mail.
- Certain variables may be very expensive to measure for all study participants.

## Missingness Mechanisms

- **MCAR:** The missing completely at random (MCAR) mechanism assumes, that the probability that a specific variable  $X$  is missing is does not depend on any other factors.
- **MAR:** The missing at random (MAR) mechanism assumes, that the probability that a specific variable  $X$  is missing does not depend on the missing value of  $X$ , but only on fully observed variables.
- **MNAR:** Otherwise data are not missing at random (MNAR).



missing\_data.png

# Fitting a linear regression model to the chain data set



```
# R performs a complete case analysis using only rows with no NAs  
lm.mod.NA <- lm(log_virus ~ age + income + healthy + mental +  
               damage + treatment, data = chain)
```

```
## Error in is.data.frame(data): object 'chain' not found
```

```
summary(lm.mod.NA)
```

```
## Error in summary(lm.mod.NA): object 'lm.mod.NA' not found
```

# Checking the model assumptions for the chain data set



```
# R performs a complete case analysis using only rows with no NAs  
par(mfrow=c(2,2))  
plot(lm.mod.NA)
```

```
## Error in plot(lm.mod.NA): object 'lm.mod.NA' not found
```



- Complete case analysis uses only those records with complete observations
- If the predictor values are missing nondifferentially, i. e. the probability that a risk factor is missing does not depend on outcome  $y$  (disease status), then complete case analysis produces valid regression coefficient estimates
- If the predictor values are missing differentially the complete case analysis can result substantial biased estimates.
- There may also be a huge loss in precision (large CIs and  $p$ -values) if a substantial part of the observations are omitted due to missing values.

## How are missing data treated?

Simple solutions that often cause bias or other problems:

- Complete data analysis = Listwise deletion
- Pairwise deletion
- Single imputation (not discussed)
- Last observation carried forward and baseline observation carried forward (not discussed)

Better (but also more complex) approaches are

- **Multiple imputation:** mainly used to handle item missings - only parts of the variables are missing
- Inverse probability weighting (not discussed)



## Iterative methods to get an imputed data set

If we have missings in more than one variable:

- **Multivariate Imputation by Chained Equations (MICE, Van Buuren and Oudshoorn (1999))**

MICE imputes each variable stochastically using a regression model conditional on all the others, iteratively cycling through all the variables that contain missing data until convergence ("Gibbs sampling"). In R we can use the mice package. The mice function assumes in default settings linearity of associations among X variables and between X and Y in the default setting, but specific forms of imputation models can be specified by the user.

- **An iterative nonparametric missing value imputation for mixed-type data** can also be done by a method called missForest which is based on a random forest approach. It can cope with high-dimensional data and may outperform MICE if complex interactions and nonlinear relations are suspected. In R we can use the missForest package.

## Iterative imputation based on RandomForest


In the R package missForest (2012) is an iterative imputation method. All imputation models are random forest models.



```
# install.packages("missForest")
library("missForest")
# Nonparametric Missing Value Imputation using Random Forest
chain.imp <- missForest(chain, variablewise = TRUE)
chain.imp$OOBError # estimated imputation error for each variable
chain.final.rf <- chain.imp$ximp
dim(chain.final.rf) # [1] 532 7
apply(chain.final.rf, 2, function(x){sum(is.na(x))})
# log_virus      age      income      healthy      mental      damage treatment
#           0           0           0           0           0           0           0
lm.mod.rf <- lm(log_virus ~ age + income + healthy +
               mental + damage + treatment,
               data = chain.final.rf)
```

```
## Error in library("missForest"): there is no package called
'missForest'
## Error in missForest(chain, variablewise = TRUE): could not find
function "missForest"
## Error in eval(expr, envir, enclos): object 'chain.imp' not found
## Error in is.data.frame(data): object 'chain.final.rf' not found
```

# Main steps used in multiple imputation with `mice` in R



# Multiple and iterative imputation with `mice` in R



```
# install.packages("mice")
library("mice")
# Multivariate Imputation by Chained Equations (MICE)
# create 4 imputed data sets via mice
# IMPUTE DATA
chain.mice <- mice(chain, m = 4, print = FALSE, seed = 2017)
# ANALYSE RESULTS WITH IMPUTED DATA
# fit for each of the 4 data sets a linear regression
fit.4.mice <- with(chain.mice, lm(log_virus ~ age + income + healthy +
                                mental + damage + treatment))

# check out the regression coefficients of the 4 fits
summary(fit.4.mice)
# POOL ANALYSIS RESULTS
# pool the results of the 4 models with Rubin's rule
lm.mod.mice <- pool(fit.4.mice)
# combined regression coefficients of the 4 fits
summary(lm.mod.mice)
```

```
## Error in library("mice"): there is no package
called 'mice'
```

```
## Error in mice(chain, m = 5, print = FALSE, seed =
2017): could not find function "mice"
```

```
## Error in with(chain.mice, lm(log_virus ~ age +
income + healthy + mental + : object 'chain.mice'
not found
```

```
## Error in pool(fit.4.mice): could not find
```

## Comparison of results: Complete case vs. missForest



```
round(summary(lm.mod.NA)$coef, 2)
```

```
## Error in summary(lm.mod.NA): object 'lm.mod.NA' not found
```

```
round(summary(lm.mod.rf)$coef, 2)
```

```
## Error in summary(lm.mod.rf): object 'lm.mod.rf' not found
```

## Comparison of results: Complete case vs. mice



```
round(summary(lm.mod.NA)$coef, 2)
```

```
## Error in summary(lm.mod.NA): object 'lm.mod.NA' not found
```

```
round(summary(lm.mod.mice)[,c(1,2,3,5)], 2)
```

```
## Error in summary(lm.mod.mice): object 'lm.mod.mice' not found
```

## Multiple Imputation - the principle idea

- $m$  imputed data sets are created instead of a single imputed data set (Rubin has shown that often  $m = 5$  to  $10$  is sufficient)
- To create the  $m$  imputed data sets take random draws from the predictive distribution from an imputation model, e.g. 5 times or do iterative imputation several times with different starting values and draw random draws.
- The variation among the  $m$  imputations reflects the uncertainty with which the missing values can be predicted from the observed data.
- With each completed data set the planned analysis is performed.
- The results from the different analyses are then combined using Rubin's rule properly taking into account the uncertainty in the imputed values.

## Summary on imputation

Missing data lead to

- Inefficient analyses of research questions
- Difficulties in interpretation when analyses differ in numbers of subjects
- Possible bias in regression coefficients
- **Imputation methods make the assumption of MAR**
- The MAR assumption is not testable, but becomes more reasonable with imputation models that include a wide range of variables.
- **Multiple imputation is recommended and superior to complete data analysis.**
- Single imputation methods are also reasonable for prediction.



## Summary for Anna Mary

- ONE continuous vs. ONE categorical variable (2 levels)
  - `t.test(...)`
- ONE continuous vs. ONE categorical variable ( $\geq 2$  levels)
  - `aov(...)`
- ONE categorical variable (2 levels) vs. ONE categorical variable (2 levels)  $\rightarrow$  2 x 2 table
  - `chisq.test(...)`
  - `fisher.test(...)`

$\rightarrow$  Regression: `lm(...)` & `glm(...)`